

Rensselaer Polytechnic Institute
CSCI-4210 – Operating Systems
David Goldschmidt, Ph.D.

Operating Systems {week 01.a}

What is an operating system?

- What is an Operating System?
 - The software interface between hardware and its users
- Operating systems:
 - Execute user and system programs
 - Manage and coordinate computer hardware
 - Serve as resource allocators
 - Are typically interrupt-driven



System software

- What is system software?
 - Computer programs that directly control the operations of the computer and its devices
- Operating systems:
 - Coordinate and orchestrate all activities of the hardware devices in a computer
 - Provide both a *Graphical User Interface (GUI)* and a *Command-Line Interface (CLI)* for its users

Operating system design goals

- From a user's perspective:
 - easy to use
 - easy to learn
 - reliable
 - safe
 - fast
 - etc.
- System goals:
 - reliability
 - flexibility
 - extensibility
 - speed(y)
 - efficiency
 - maintainability
 - etc.



Operating system services (i)

- An operating system provides services:
 - Program execution
 - Load programs into memory, run/suspend/halt programs, handle/display errors
 - I/O operations
 - Seamlessly interact with I/O devices, including disks, networks connection, etc.
 - Filesystem manipulation
 - Read/write/traverse filesystem directories, read/write files, enforce permissions, search for files

Operating system services (ii)

- Other operating system services:
 - Inter-Process Communications (IPC)
 - Processes exchange information via shared memory, message passing, sockets, pipes, files, etc.
 - Often spans multiple computers and networks
 - Error detection and recovery
 - Detect errors in CPU, memory, I/O devices, processes, network connections, etc.
 - Recover from errors gracefully, ensuring correct and consistent operations




Batch jobs

- A **job** is a unit of work submitted by a user to the operating system
- Jobs typically consist of:
 - a program either in a source language or in "executable" binary form
 - input data used by the program when it executes




Multiprogramming (i)



Multiprogramming (ii)

- In multiprogramming, several jobs reside in memory simultaneously
 - CPU use is shared and managed by the operating system




Multiprogramming (iii)

- Multiprogramming provides efficient use of the computer (CPU) and its resources (I/O)
 - One user cannot keep the CPU and I/O devices busy at all times
 - Multiprogramming organizes jobs such that the CPU always has exactly one to execute

Multiprogramming (iv)

- Computer is often idle – why?
 - CPU and hardware *significantly* faster than I/O
 - When a user or process is blocked waiting for I/O, the operating system switches to another job
 - A subset of jobs is stored in memory, awaiting CPU or I/O



Timesharing (a.k.a. multitasking)

- To ensure fairness, use a *timesharing scheme* in which the CPU cycles through all jobs
 - Each job is given a fixed amount of CPU time
 - Switching from one running job (or process) to another is called a *context switch*
 - A process may relinquish its time if blocked on an I/O request




We interrupt this program...

- Software instructions are executed by a Central Processing Unit (CPU)
 - An external hardware event triggers an interrupt by signaling the CPU
 - e.g. mouse movement, keyboard event
 - Software triggers an interrupt by executing a system call
 - e.g. disk read, graphics output, printer output




Interrupt mechanism


- Interrupts are handled much like calling a function in a programming language




Typical interrupt timeline



Synchronous and Asynchronous I/O




Hierarchical storage architecture



Caching (i)

- Caching is a technique in which data is temporarily stored in a smaller and faster memory component
 - Why implement caching in an operating system?



Caching (ii)

- A key goal in operating system design is achieving *fast and efficient performance*

Level	1	2	3	4
Name	registers	cache	main memory	disk storage
Typical size	< 1 KB	> 16 MB	> 16 GB	> 100 GB
Implementation technology	custom memory with multiple ports, CMOS	on-chip or off-chip CMOS DRAM	CMOS DRAM	magnetic disk
Access time (ns)	0.25 – 0.5	0.5 – 25	80 – 250	5,000,000
Bandwidth (MB/sec)	20,000 – 100,000	5000 – 10,000	1000 – 5000	20 – 150
Managed by	compiler	hardware	operating system	operating system
Backed by	cache	main memory	disk	CD or tape

Caching (iii)

- What's the caching algorithm?
 - When the operating system attempts to read from memory, check to see if the requested data is already in the cache
 - If it is, data is read from the cache (*fast!*)
 - If not, data is copied from memory to the cache (*maybe next time...*)



Principle of locality

- When a running program reads from memory location X, the *principle of locality* predicts that the next memory location requested will be near X



- Store *pages* of data in a cache, where each page is typically the same size (e.g. 64KB)

Programming exercise



- Implement a program to simulate caching:
 - Write a function called `calculateAnswer()` that takes integer `n` as input and calculates (and returns) the sum $(1 + 2 + \dots + n)$
 - Pretend this method is computationally costly!
 - Initially, the cache is empty
 - Ask the user to input a number in range 1..100
 - If the answer is not in the cache, call `calculateAnswer()` and display the resulting sum; store the result in the cache
 - If the answer is in the cache, simply display the answer